

CSC 481: First Order Logic

c- Pragmatics and conclusion

Rodrigo Canaan
Assistant Professor
Computer Science Department
Cal Poly, San Luis Obispo
rcanaan@calpoly.edu

The basics


A language is composed of:

- **Syntax:** what symbols may be used and what combinations of symbols are well-formed
 - “I drove to work today” is a well-formed English sentence
 - “I work to today drove” is not
- **Semantics:** what each sentence means
 - Sentences are used to convey that the world is one way and not another
- **Pragmatics:** what the language is used for
 - The sentence “fire!” suggests a different response in a crowded theater or in a shooting range

The basics

A language is composed of:

- **Syntax:** what symbols may be used and what combinations of symbols are well-formed
 - “I drove to work today” is a well-formed English sentence
 - “I work to today drove” is not
- **Semantics:** what each sentence means
 - Sentences are used to convey that the world is one way and not another
- **Pragmatics:** what the language is used for
 - The sentence “fire!” suggests a different response in a crowded theater or in a shooting range



Last class

Satisfaction and Models

Assume that t_1, \dots, t_n are terms, P is a predicate of arity n , α and β are formulas, and x is a variable.

1. $\mathfrak{S}, \mu \models P(t_1, \dots, t_n)$ iff $\langle d_1, \dots, d_n \rangle \in \mathcal{P}$, where $\mathcal{P} = \mathcal{I}[P]$, and $d_i = \|t_i\|_{\mathfrak{S}, \mu}$;
2. $\mathfrak{S}, \mu \models t_1 = t_2$ iff $\|t_1\|_{\mathfrak{S}, \mu}$ and $\|t_2\|_{\mathfrak{S}, \mu}$ are the same element of \mathcal{D} ;
3. $\mathfrak{S}, \mu \models \neg\alpha$ iff it is not the case that $\mathfrak{S}, \mu \models \alpha$;
4. $\mathfrak{S}, \mu \models (\alpha \wedge \beta)$ iff $\mathfrak{S}, \mu \models \alpha$ and $\mathfrak{S}, \mu \models \beta$;
5. $\mathfrak{S}, \mu \models (\alpha \vee \beta)$ iff $\mathfrak{S}, \mu \models \alpha$ or $\mathfrak{S}, \mu \models \beta$ (or both);
6. $\mathfrak{S}, \mu \models \exists x.\alpha$ iff $\mathfrak{S}, \mu' \models \alpha$, for some variable assignment μ' that differs from μ on at most x ;
7. $\mathfrak{S}, \mu \models \forall x.\alpha$ iff $\mathfrak{S}, \mu' \models \alpha$, for every variable assignment μ' that differs from μ on at most x .


1. $\text{married}(\text{'Ryan'}, \text{'Blake'})$
2. $\text{spouse}(\text{'Ryan'}) = \text{'Blake'}$
3. $\neg \text{dog}(\text{'Garfield'})$
4. $\text{dog}(\text{'Toto'}) \wedge \text{married}(\text{'Ryan'}, \text{'Blake'})$
5. $\text{dog}(\text{'Garfield'}) \vee \text{dog}(\text{'Toto'})$
6. $\exists x \text{ married}(x, \text{'Blake'})$
7. $\forall x \text{ dog}(x) \rightarrow \text{mammal}(x)$

We say $\mathfrak{S} \models \alpha$ or α is satisfied in \mathfrak{S} or (conversely) \mathfrak{S} satisfies α

The basics

A language is composed of:

- **Syntax:** what symbols may be used and what combinations of symbols are well-formed
 - “I drove to work today” is a well-formed English sentence
 - “I work to today drove” is not
- **Semantics:** what each sentence means
 - Sentences are used to convey that the world is one way and not another
- **Pragmatics:** what the language is used for
 - The sentence “fire!” suggests a different response in a crowded theater or in a shooting range



Today

Pragmatics

Our goal

What do we want a KB for?

- Given a set of logical sentences assumed to represent relevant knowledge of the world (a knowledge base, or KB), we may want to:
 - Use the KB for querying: does some other fact follow from our KB?
 - Example: given the facts below, does it follow that Bob is the killer?
 - $killer(Alice) \vee killer(Bob)$
 - $hasAlibi(Alice)$
 - $\forall x : hasAlibi(x) \rightarrow \neg killer(x)$

Our goal

What do we want a KB for?

- Given a set of logical sentences assumed to represent relevant knowledge of the world (a knowledge base, or KB), we may want to:
 - Use the KB for as basis for an agent: what is the best action to do?
 - Example: given a certain Minesweeper board, which hidden tile should we explore next?
 - This can be achieved in several ways:
 - we could query for *safe(tile)* and have our agent click that tile if so
 - or we could define a predicate *bestTileToExplore/1* and some rule to produce it.

Our goal

What do we want a KB for?

- In both cases, we want the retrieve answer/action to “logically follow” from the KB
 - But what does this mean?

Logical Consequence

Intuition

- Given an interpretation, we can determine the truth value of any sentence
- However, our program has no idea what our intended interpretation is
- We would still like our program to be able to derive conclusions that agree with our interpretation

Logical Consequence

Intuition (cont.)

- Given the sentences:
 - $S1$ = “Fido is a dog”
 - $S2$ = “All dogs are mammals”
 - α = “Fido is a mammal”
- Each of these sentences may or may not be true for a given interpretation \mathfrak{I}
- What we can do is say: if \mathfrak{I} happens to satisfy both $S1$ and $S2$, then it must satisfy α

Logical Consequence

Valid and unsatisfiable sentences

- Note also that the meaning of some sentences does not depend on any interpretation. For example, if α is any sentence:
 - $\alpha \vee \neg\alpha$ is always true (α is **valid**)
 - $\alpha \wedge \neg\alpha$ is always false (α is **unsatisfiable**)
- Valid sentences are also called tautologies
- Invalid sentences are also called contradictions


Logical Consequence

Definition

- Given a (finite) set of sentences $S = \{s_1, s_2, s_3, \dots, s_n\}$ and a sentence α , the following statements are equivalent:
 - $S \models \alpha$ (read “**S entails α** ”)
 - For every \mathfrak{I} , if $\mathfrak{I} \models S$, then $\mathfrak{I} \models \alpha$
 - The set of interpretations that satisfy S is a subset of the set of interpretations that satisfy α
 - There is no \mathfrak{I} such that $\mathfrak{I} \models S \cup \{\neg\alpha\}$
 - $S \cup \{\neg\alpha\}$ is unsatisfiable
 - $(s_1 \wedge s_2 \wedge s_3 \wedge \dots \wedge s_n) \rightarrow \alpha$ is a tautology
 - $(s_1 \wedge s_2 \wedge s_3 \wedge \dots \wedge s_n) \wedge \neg\alpha$ is a contradiction

Models

- Logicians typically think in terms of **models**, which are formally structured worlds with respect to which truth can be evaluated
 - We say m is a model of a sentence α if α is true in m
 - $M(\alpha)$ is the set of all models of α
 - ❖ Then $KB \models \alpha$ iff $M(KB) \subseteq M(\alpha)$
- E.g.
- ❖ KB = Giants won and Reds won
 - ❖ α = Giants won



Models

- Logicians typically think in terms of **models**, which are formally structured worlds with respect to which truth can be evaluated

- We say m is a model of a sentence α if α is true in m


- $M(\alpha)$ is the set of all models of α

- ❖ Then $KB \models \alpha$ iff $M(KB) \subseteq M(\alpha)$

E.g.

- ❖ KB = Giants won and Reds won

- ❖ α = Giants won



Note: KB is a stronger assertion than α


Example: block world

- An interpretation corresponds to an arrangement of blocks
 - Domain: a set of blocks
 - Predicates: blue/1, onTopOf/2

Example: block world

- An interpretation corresponds to an arrangement of blocks
 - Domain: a set of blocks
 - Predicates: blue, onTopOf


A possible interpretation




Example: block world

- An interpretation corresponds to an arrangement of blocks
 - Domain: a set of blocks
 - Predicates: blue, onTopOf

A possible interpretation




Another interpretation




Example: block world

- An interpretation corresponds to an arrangement of blocks
 - Domain: a set of blocks
 - Predicates: blue, onTopOf


A possible interpretation



Another interpretation




Some interpretations may not even make physical sense




Example: block world

- What if the agent does not see the color of all blocks?
- Is there a blue block on top of a non blue block?



Example: block world

- Only 2 interpretations satisfy what the agent knows
- In both of them, there is a blue block on top of a non-blue, so yes



Deductive inference

- It's easy to explain the conclusion from the previous slide in English
- We want to do it mechanically
- Idea:
 - Express agent knowledge as symbols in a KB
 - Express the goal as the symbols of a “query” sentence α
 - Manipulate the KB symbols until we produce the symbols for “ α is true” or “ α is false”


Deductive inference

- If a given algorithm i succeeds in proving α from the KB, we say it derives α
 - $\text{KB} \vdash_i \alpha$
- We want $\text{KB} \vdash_i \alpha$ if and only if $\text{KB} \models \alpha$
 - If an algorithm only derives sentences that are actually entailed, we say it is **sound**
 - If it always produces an answer, we say it is **complete**
- Unfortunately, no sound and complete algorithm for checking the entailment of sentences in First-Order Logic can exist

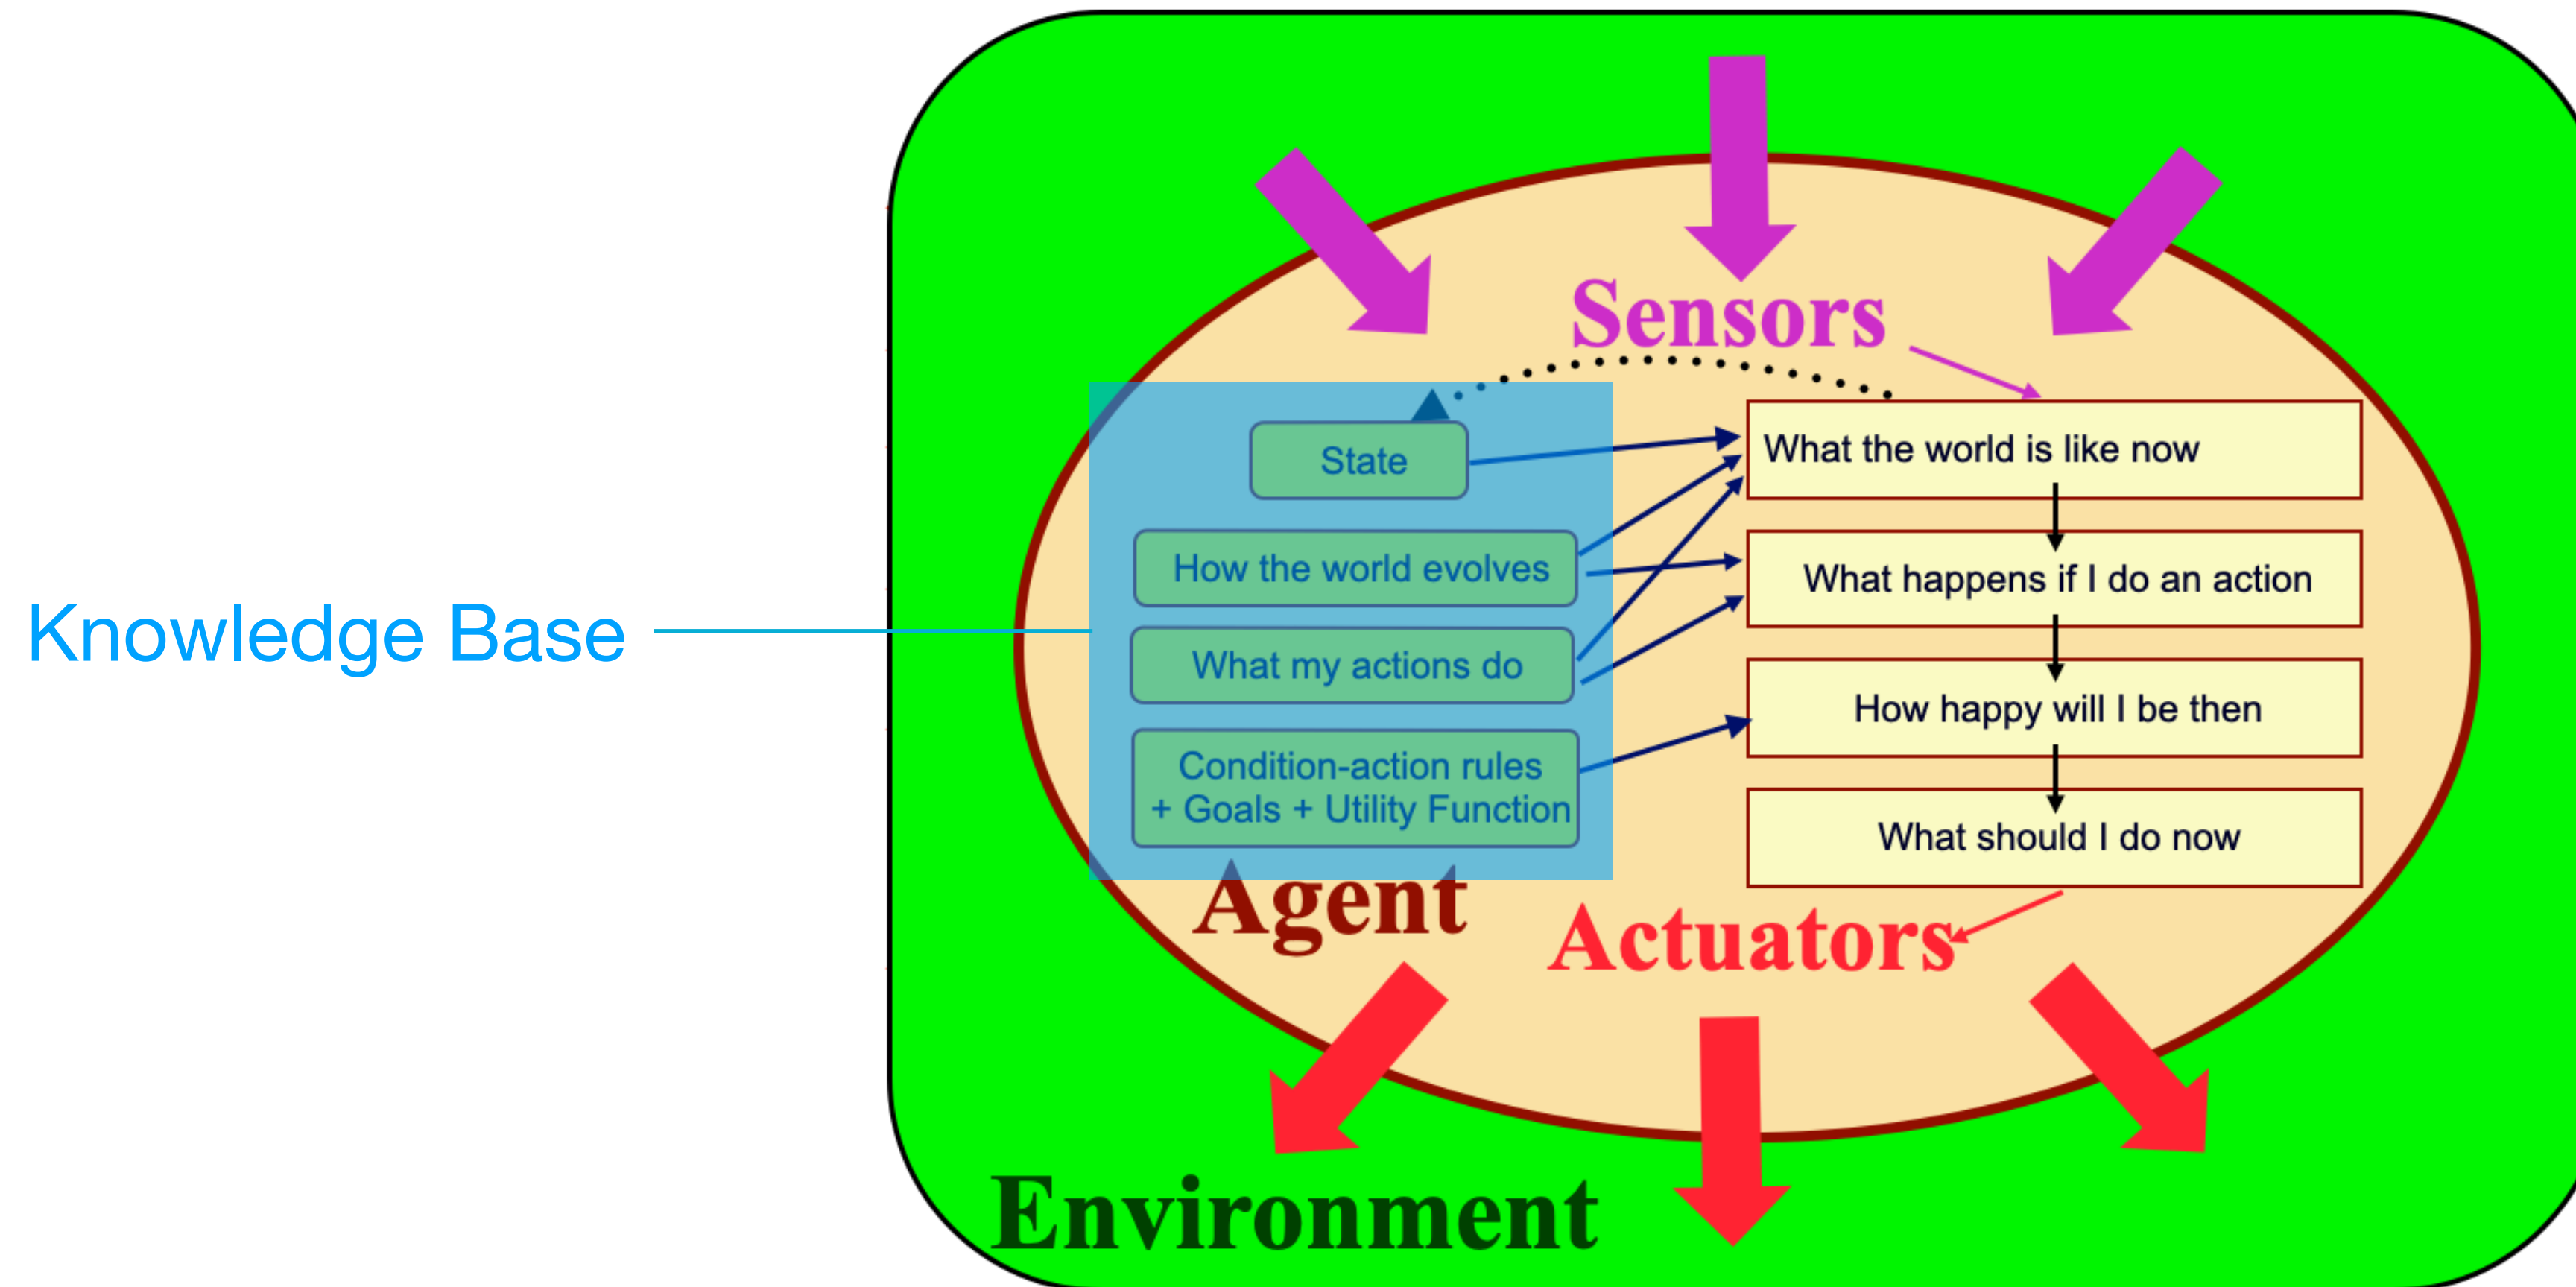
A note on agents

- Strictly speaking, logic tells us what knowledge can be inferred from a KB (that is, producing new knowledge from existing knowledge)
- However, we are often interested in *what is the right thing to do*.
- An **agent** is an entity (machine or human) that interacts with an environment trying to achieve a goal
- To bridge the gap between knowledge and action, we many want to check whether the KB entails propositions of the form “The correct thing to do is X”
- An agent that acts according to this principle is often called a **knowledge-based agent**

Example of Agent Diagram




Example of Agent Diagram



Knowledge-Based agents

Inference rules such as:

- $A \vee B, \neg A \vdash B$



<http://aima.eecs.berkeley.edu/slides-ppt/>

“Facts” such as

- Fido is a dog
- There is a 2 in the top-left square

“Rules” such as

- Every dog is a mammal
- If two squares share a column, they must have different numbers

Knowledge-Based agents pseudocode

```
function KB-AGENT(percept) returns an action
  static: KB, a knowledge base
          t, a counter, initially 0, indicating time

  TELL(KB, MAKE-PERCEPT-SENTENCE(percept, t))
  action ← ASK(KB, MAKE-ACTION-QUERY(t))
  TELL(KB, MAKE-ACTION-SENTENCE(action, t))
  t ← t + 1
  return action
```


Knowledge-Based agents pseudocode

```
function KB-AGENT(percept) returns an action
  static: KB, a knowledge base
          t, a counter, initially 0, indicating time

  TELL(KB, MAKE-PERCEPT-SENTENCE(percept, t))
  action ← ASK(KB, MAKE-ACTION-QUERY(t))
  TELL(KB, MAKE-ACTION-SENTENCE(action, t))
  t ← t + 1
  return action
```


Building a KB

- The ALMA book suggests the following steps for knowledge engineering projects (ch. 8.4.1):

1. Identify the questions

2. Assemble the relevant knowledge

3. Decide on a vocabulary (predicates, functions, constants)

4. Encode general knowledge about the domain

5. Encode a description of the project instance

6. Pose queries to the inference procedure and get answers

7. Debug and evaluate the KB

Next class

- We will apply those steps to some more interesting problems
- Start reading chapter 3 of the book, entry ticket will be due Thursday
- I will assign a quiz based on chapter 2 in the next few days.